AI-Powered Matchmaking Tool for Proof of Talk

A) Executive Summary

Proof of Talk wins on conversation quality, not networking volume. In a curated room of senior decision-makers, weak introductions are expensive: they waste scarce meeting slots and reduce trust in the event experience. This proposal delivers a quality-first matchmaking system that recommends who to meet before and during the summit, with clear reasons people can trust. It combines hard constraints, strategic-fit scoring, and limited discovery to avoid echo chambers. It is consent-first: registration is required, enrichment is opt-in. It also surfaces non-obvious opportunities, including pair and triad pathways (for example, founder -> GP -> LP). Output is intentionally constrained to high-confidence results only (score > 65, target 3-7 recommendations when available). The MVP is realistic to ship in weeks and improves during the event through live feedback.

B) Problem Framing

Why matchmaking is uniquely hard at Proof of Talk

• Attendees are senior and time-constrained.

• Goals are asymmetric: fundraising, partnerships, regulation, hiring, learning.

• Job titles do not reliably reflect current priorities.

• Bad matches damage trust in a high-curation environment.

Why simple filtering fails

• Keyword overlap creates generic matches.

• Similarity-only logic misses complementary value.

• Static directories cannot adapt to in-event context changes.

What success means

• Fewer, higher-confidence introductions.

• Higher “useful meeting” and “follow-up planned” rates.

• Better organizer outcomes: curated tables, stronger intros, measurable business progress.

C) Users & Personas

1) Institutional Investor (Partner/MD)

• Goals: differentiated deal flow, thesis validation.

• Constraints: very limited slots, low tolerance for generic outreach.

• Good match: founder/executive with clear fit, credibility, and timing.

2) Founder/CEO

• Goals: capital, strategic distribution, enterprise partnerships.

• Constraints: high opportunity cost per meeting.

• Good match: decision-maker with direct relevance and a realistic next step.

3) Bank / Institutional Innovation Executive

• Goals: compliant digital-asset adoption and partner selection.

• Constraints: regulatory and reputational risk.

• Good match: enterprise-ready team with strong compliance posture.

4) Infrastructure / Protocol CTO

• Goals: technical partnerships, design partners, selective hiring.

• Constraints: avoids low-depth, non-technical meetings.

• Good match: counterpart with architecture-level complementarity.

D) Data Strategy

Sources, value, and risk

1. Registration form (mandatory)

• Value: structured baseline and explicit intent.

• Risk: often shallow.

2. LinkedIn (optional opt-in)

• Value: expertise and trajectory context.

• Risk: privacy sensitivity; stale fields.

3. Company website (optional)

• Value: company strategy and positioning.

• Risk: marketing bias.

4. Public content (optional)

• Value: current thinking and active themes.

• Risk: uneven availability.

5. Past Proof of Talk data (if available)

• Value: strongest context signal for quality.

• Risk: can reinforce old network patterns.

Mandatory vs optional

• Mandatory: profile basics, goal, availability, exclusions.

• Optional: LinkedIn, website, public content.

Consent and noisy data handling

• Source-level opt-in.

• Plain-language data use disclosure.

• Attendee controls for edits/exclusions.

• Confidence labels for inferred fields.

• Conservative ranking for low-confidence profiles.

Cold-start strategy

• 2-minute intent form (seek, offer, constraints).

• Rules-first matching plus organizer curation fallback.

E) Profile Building Logic

Each attendee is represented by four blocks:

• Identity context (role, authority, language, sector)

• Current focus (what they are actively doing now)

• Seeks (what they want from this event)

• Offers (what value they can provide)

Text signals are interpreted semantically, but structured fields remain the control layer.

Conflict policy:

• Explicit attendee input overrides inferred signals.

• Newer signals outrank older ones.

• Low-confidence inferences are down-weighted.

F) Matching Logic (Core)

1) Hard constraints

• Language compatibility

• Availability overlap

• Exclusions and organizer policy rules

2) Soft scoring

• Objective relevance

• Complementarity (investor-founder, bank-infra, policy-operator)

• Decision-level fit

• Asymmetric value exchange

• Feedback prior from observed outcomes

3) Exploration / diversity injection

• One limited discovery slot for adjacent high-upside opportunities.

• Discovery still must pass quality threshold.

4) Explainability

Each recommendation includes clear reasons, e.g.:

• “Your investment objective aligns with this founder’s active fundraising and institutional traction.”

• “You both have overlapping availability and compatible strategic priorities.”

5) Strategic scenario engine (non-obvious opportunities)

• Pair: Kenji’s compliance-focused L2 -> Elena’s bank deployment need.

• Triad: Marcus (founder) -> Aisha (GP) -> Sarah (LP context bridge).

G) Output & UX Design

Before event (web-first)

• Show only high-confidence recommendations (`score > 65`).

• Target 3-7 recommendations when enough qualify.

• Match cards include reasons and actions: Request Intro / Save / Not Relevant.

During event

• Time-aware nudges based on free windows.

• Concierge escalation on mutual accept.

After event

• Outcome capture: met, follow-up planned, declined.

• Feedback updates future ranking.

Example copy:

“Recommended: Elena Rossi (European Bank). Why: your objective is institutional allocation, and Elena is launching compliant custody while seeking tokenization partners. Shared availability: Day 1 PM.”

H) Learning & Feedback Loop

• Explicit: rating, outcome, comment.

• Implicit: accept/decline, meeting completion, follow-up behavior.

• Re-rank during event without heavyweight retraining.

• Repeated weak outcomes are penalized or removed.

I) System Design (Logical, Not Over-Engineered)

Core components

• Data ingestion

• Profile builder

• Matching service

• Explanation layer

• Scenario engine (pair/triad)

• Notification layer

• Organizer dashboard

High-level schema

• `attendees`

• `matches`

• `feedback`

• `signals` (optional enrichment traces)

APIs

• `GET /v1/matches/{attendee_id}`

• `POST /v1/feedback`

• `GET /v1/organizer/metrics`

• `GET /v1/scenarios`

```
[Ingestion] -> [Profile Builder] -> [Matching + Scenario Engine] -> [Explanations]
                                                |                          |
                                                v                          v
                                        [Organizer Tools]          [Attendee UI]
```

J) Organizer Perspective

Organizers can use this system to:

• prioritize concierge introductions,

• design better seating and roundtables,

• detect high-value attendees with weak coverage,

• act on pair/triad scenario opportunities.

This strengthens Proof of Talk’s perceived and measurable event value.

K) Risks, Ethics & Privacy

• Consent-first enrichment; no hidden scraping.

• Transparent “why this match” explanations.

• Anti-spam controls and moderation paths.

• Bias control through exploration and organizer oversight.

• Deliberate non-automation: final intro approval remains human.

L) MVP Scope & Roadmap

Level 1 (concept)

• End-to-end design, data strategy, matching logic, UX flow, risks, and roadmap.

Level 2 (delivered)

• Clickable wireframe.

• Working POC with 12 fictional attendees.

• Generated sample input/output, including strategic scenarios.

• Working web app with attendee workflow + organizer workspace.

• Organizer attendee input form and CSV export for match recommendations.

• External data retrieval from one source (company website enrichment endpoint).

• Explicit LinkedIn opt-in enrichment flow (checkbox + profile URL; no hidden scraping).

• Production-grade presentation layer improvements:

• Proof of Talk branding integration (local static logo),

• browser favicon support,

• premium UI interactions with lightweight motion,

• responsive spacing stabilization across mobile/tablet/laptop/desktop/large monitor viewports.

90-day roadmap

1. Days 1-30: ingestion, profile builder, hard constraints, quality-gated recommendations.
2. Days 31-60: scoring refinement, explanation quality, feedback loop, organizer metrics.
3. Days 61-90: in-event nudges, scenario tuning, stronger organizer curation workflows.

Technical Stack

• Web-first UI

• FastAPI backend

• SQLite for MVP (Postgres in production)

• Rule-based scoring + semantic text interpretation + feedback priors

• HTTP enrichment adapter for external source retrieval (company website metadata/content summary)

• Security and runtime hardening:

• RBAC, CSRF, lockout/rate limiting, audit logs, SSRF controls, secure headers.

Why this differs from Grip / Brella / Swapcard

• Strict quality gating, not high-volume networking.

• Complementary and strategic-chain logic, not tag similarity only.

• Explainability and organizer-in-the-loop controls as core features.

How This Scores Against Evaluation Criteria

• Problem Understanding: high-stakes framing, edge cases, quality KPIs.

• System Design: practical architecture, consent model, organizer controls.

• AI & Matching Logic: constraints + scoring + explainability + scenarios.

• Communication: concise, executive-readable, concrete examples.

• Ambition & Execution: Level 1 + Level 2 artifacts + runnable prototype with authenticated organizer tooling, CSV export, external source enrichment, and cross-device UI reliability improvements.

“Things to Think About” — Direct Answers

1. Cold start: require mini-intent form, rank conservatively, use organizer fallback.
2. Similar vs complementary: prioritize complementary; keep one controlled peer-similarity slot.
3. Quality vs quantity: enforce `score &gt; 65`, target 3-7 recommendations.
4. Privacy/consent: source-level opt-in, transparent use, attendee edit controls.
5. Organizer value: seating, roundtables, concierge intros, coverage monitoring.
6. Feedback loops: explicit + implicit signals drive reranking.
7. Existing tools: differentiate through quality gating, explainability, and triad scenarios.

Level 2 Package

• Wireframe: `/docs/level2/wireframe-clickable.html`

• POC note: `/docs/level2/Proof_of_Concept.md`

• Input: `/docs/level2/sample_input_12_attendees.json`

• Match output: `/docs/level2/sample_output_matches.json`

• Scenario output: `/docs/level2/sample_output_scenarios.json`

• Working app endpoints include:

• `GET /v1/matches/{attendee_id}`

• `POST /v1/enrich/company`

• `POST /v1/enrich/linkedin`

• `GET /organizer/export/matches.csv`

• `GET /v1/scenarios`

• `GET /favicon.ico` (brand icon support)

Submission Email Paragraph

I am submitting my case study for an AI-powered matchmaking tool for Proof of Talk. The proposal is quality-first and implementation-focused: consented profile enrichment, explainable matching, strict confidence gating, and feedback-driven adaptation. I also included a Level-2 wireframe and working proof of concept with sample outputs, including non-obvious pair and triad strategic opportunities.

Evaluation Checklist

| Criteria | Coverage |

|---|---|

| Problem Understanding (30%) | High-stakes framing, edge cases, quality KPIs |

| System Design (25%) | Practical architecture, consent model, organizer controls |

| AI & Matching Logic (25%) | Hard constraints, soft scoring, explainability, scenarios |

| Communication (10%) | Concise, structured, executive-readable narrative |

| Ambition & Execution (10%) | Level 1 + Level 2 artifacts + runnable prototype |