Level 3 Prototype Status and Quality Policy

This working prototype is configured with a strict quality-over-quantity policy:

- Show only matches with score > 65.

- Return 3-7 matches when enough candidates pass threshold.

- If fewer than 3 pass, the system shows only qualified matches and flags organizer curation fallback.

- Strategic scenario engine surfaces pair and triad opportunities for concierge intros.

Current implementation status

- Running web app with attendee view and organizer snapshot.

- Role-based authentication is active (organizer and attendee access paths).

- Per-attendee credential model is active (attendee123-<id> for seeded demo accounts).

- API endpoints are role-protected with session auth.

- Audit logs are captured for login, attendee changes, enrichment, intros, feedback, and exports.

- CSRF, rate-limiting, lockout, SSRF protections, and security headers are implemented.

- Match explanations in plain English.

- Feedback submission updates future ranking.

- Generated sample input/output artifacts for 12 fictional attendees.

- Double-opt-in intro workflow is active.

- Organizer can input attendees and export CSV recommendations.

- Organizer can bulk import attendee lists via CSV/JSON with row-level validation and partial-success reporting.

- Organizer can delete attendees with typed-name confirmation; related attendee-scoped records are removed in one action.

- External source enrichment is available via company website endpoint.

- LinkedIn enrichment is available with explicit attendee opt-in (checkbox + profile URL).

- Containerized deployment assets added (Dockerfile, docker-compose.yml, .env.example).

- Scalability guards added for 2,500-attendee operation:

- bounded scenario generation,

- paginated organizer/overview attendee lists,

- aggregated feedback scoring query,

- batched candidate hydration for match responses.

- UI/UX hardening is implemented:

- Proof of Talk logo served locally under /static/brand/proof-of-talk-logo.svg,

- favicon and /favicon.ico route support,

- premium motion with reduced-motion safe fallback,

- reveal fallback logic to prevent hidden sections after resize/orientation changes,

- spacing fixes for Curated Directory and Strategic Scenarios text cards (more generous card padding + improved line-height for readability).
- Cross-device verification executed with Playwright across:
- iPhone, Android, tablet, laptop, desktop, and big monitor viewports.

Endpoints

- GET /v1/matches/{attendee_id}
- POST /v1/feedback
- POST /v1/intros
- POST /v1/intros/{intro_id}
- POST /v1/enrich/company
- POST /v1/enrich/linkedin
- GET /v1/organizer/metrics
- GET /v1/scenarios
- GET /v1/scenarios?attendee_id=<id>
- GET /health
- GET /favicon.ico

UI routes

- GET /login
- GET / (role-aware landing)
- GET /attendees/{id} (attendee workspace, role-scoped)
- GET /organizer (organizer workspace)
- GET /organizer/attendees/template.csv (import template)
- POST /organizer/attendees/import (organizer-only bulk import)
- POST /organizer/attendees/{id}/delete (organizer-only, CSRF + confirmation gated)

Demo Credentials

- Organizer email: organizer@pot.local
- Organizer password: organizer123
- Attendee login ID example: 1
- Attendee passcode pattern: attendee123-<attendee_id> (example: attendee123-1)

Why this improves quality

- Prevents weak "just in case" recommendations.
- Preserves attendee trust in curated intros.
- Keeps meeting calendars focused on high-confidence opportunities.
- Gives organizers explicit pair/triad pathways for high-leverage intros.

Security reference document:

- /docs/security/Security_Hardening_Plan_and_Implementation.md

Current deployed URL:

- https://ai-matchmaking-pot.vercel.app (https://ai-matchmaking-pot.vercel.app)

Primary wireframe (Figma):

- Proof of Talk Pitch Board (https://www.figma.com/online-whiteboard/create-diagram/f832226c-9b28-4c38-be 43-8b4ada8f4d64?utm_source=other&utm_content=edit_in_figjam&oai_id=&request_id=55701c3a-ed32-40 8d-90b0-bb827b8e79fd)