

Level 3 Prototype Status and Quality Policy

This working prototype is configured with a strict quality-over-quantity policy:

- Show only matches with score > 65.
- Return 3-7 matches when enough candidates pass threshold.
- If fewer than 3 pass, the system shows only qualified matches and flags organizer curation fallback.
- Strategic scenario engine surfaces pair and triad opportunities for concierge intros.

Current implementation status

- Running web app with attendee view and organizer snapshot.
- Role-based authentication is active (organizer and attendee access paths).
- Per-attendee credential model is active (attendee123-<id> for seeded demo accounts).
- API endpoints are role-protected with session auth.
- Audit logs are captured for login, attendee changes, enrichment, intros, feedback, and exports.
- CSRF, rate-limiting, lockout, SSRF protections, and security headers are implemented.
- Match explanations in plain English.
- Feedback submission updates future ranking.
- Generated sample input/output artifacts for 12 fictional attendees.
- Double-opt-in intro workflow is active.
- Organizer can input attendees and export CSV recommendations.
- Organizer can bulk import attendee lists via CSV/JSON with row-level validation and partial-success reporting.
- Organizer can delete attendees with typed-name confirmation; related attendee-scoped records are removed in one action.
- External source enrichment is available via company website endpoint.
- LinkedIn enrichment is available with explicit attendee opt-in (checkbox + profile URL).
- Containerized deployment assets added (Dockerfile, docker-compose.yml, .env.example).
- Scalability guards added for 2,500-attendee operation:
 - bounded scenario generation,
 - paginated organizer/overview attendee lists,
 - aggregated feedback scoring query,
 - batched candidate hydration for match responses.
- UI/UX hardening is implemented:
 - Proof of Talk logo served locally under /static/brand/proof-of-talk-logo.svg,
 - favicon and /favicon.ico route support,
 - premium motion with reduced-motion safe fallback,
 - reveal fallback logic to prevent hidden sections after resize/orientation changes,
- spacing fixes for Curated Directory and Strategic Scenarios text cards (more generous card padding + improved line-height for readability).
- Cross-device verification executed with Playwright across:
 - iPhone, Android, tablet, laptop, desktop, and big monitor viewports.

Endpoints

- GET /v1/matches/{attendee_id}
- POST /v1/feedback
- POST /v1/intros
- POST /v1/intros/{intro_id}

- POST /v1/enrich/company
- POST /v1/enrich/linkedin
- GET /v1/organizer/metrics
- GET /v1/scenarios
- GET /v1/scenarios?attendee_id=<id>
- GET /health
- GET /favicon.ico

UI routes

- GET /login
- GET / (role-aware landing)
- GET /attendees/{id} (attendee workspace, role-scoped)
- GET /organizer (organizer workspace)
- GET /organizer/attendees/template.csv (import template)
- POST /organizer/attendees/import (organizer-only bulk import)
- POST /organizer/attendees/{id}/delete (organizer-only, CSRF + confirmation gated)

Demo Credentials

- Organizer email: organizer@pot.local
- Organizer password: organizer123
- Attendee login ID example: 1
- Attendee passcode pattern: attendee123-<attendee_id> (example: attendee123-1)

Why this improves quality

- Prevents weak “just in case” recommendations.
- Preserves attendee trust in curated intros.
- Keeps meeting calendars focused on high-confidence opportunities.
- Gives organizers explicit pair/triad pathways for high-leverage intros.

Security reference document:

- /docs/security/Security_Hardening_Plan_and_Implementation.md

Current deployed URL:

- <https://ai-matchmaking-pot.vercel.app>

Primary wireframe (Figma):

- Proof of Talk Pitch Board